

The Design and Analysis of Algorithms

Lecture 6 Greedy Algorithms I

Zhenbo Wang

Department of Mathematical Sciences, Tsinghua University



Content

Interval Scheduling


Scheduling to Minimize Lateness

Dijkstra's Algorithm



Interval Scheduling

- Job j starts at s_j and finishes at f_j .
- Two jobs *compatible* if they don't overlap.
- Goal: find maximum subset of mutually compatible jobs.






Interval Scheduling: Greedy Algorithms

- *Greedy template.* Consider jobs in some natural order.
 - Take each job provided it's compatible with the ones already taken.
- (a) [Earliest start time] Consider jobs in ascending order of s_j .
- (b) [Shortest interval] Consider jobs in ascending order of $f_j - s_j$.
- (c) [Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .
- (d) [Earliest finish time] Consider jobs in ascending order of f_j .



Counterexamples for Greedy (a), (b) and (c)



Earliest-Finish-Time-First Algorithm

EARLIEST – FINISH – TIME – FIRST($n, s_1, \dots, s_n, f_1, \dots, f_n$)

```
1: SORT jobs by finish time so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .
2:  $A \leftarrow \emptyset$ 
3: for  $j = 1$  to  $n$  do
4:   if job  $j$  is compatible with  $A$  then
5:      $A \leftarrow A \cup \{j\}$ .
6:   end if
7: end for
8: return  $A$ .
```

- *Proposition.* Can implement earliest-finish-time first in $O(n \log n)$ time.

Keep track of job j^* that was added last to A .

Job k is compatible with A iff $s_k \geq f_{j^*}$.

Sorting by finish time takes $O(n \log n)$ time.



Analysis of Earliest-Finish-Time-First Algorithm

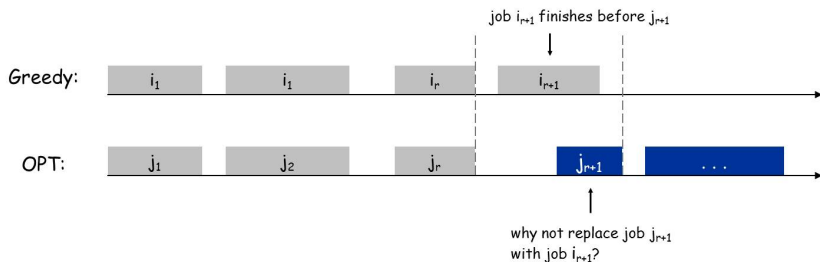
Theorem 1

The earliest-finish-time-first algorithm is optimal.

Pf. [by contradiction]

Assume greedy is not optimal. Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in the optimal solution with $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



Scheduling to Minimize Lateness

- Minimizing lateness problem.

Single resource processes one job at a time.


Job j requires t_j units of processing time and is due at time d_j .

If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.

Lateness: $l_j = \max\{0, f_j - d_j\}$.

Goal: Schedule all jobs to minimize *maximum* lateness $L = \max_j l_j$.

	1	2	3	4	5	6
t_j	3	2	1	4	3	2
d_j	6	8	9	9	14	15



Minimizing Lateness: Greedy Algorithms

- *Greedy template.* Schedule jobs according to some natural order.

[Shortest processing time first] Schedule jobs in ascending order of processing time t_j .

[Smallest slack] Schedule jobs in ascending order of slack $d_j - t_j$.

[Earliest deadline first] Schedule jobs in ascending order of deadline d_j .

	1	2
t_j	1	10
d_j	100	10

counterexample

	1	2
t_j	1	10
d_j	2	10


counterexample



Minimizing Lateness: Earliest Deadline First

EARLIEST – DEADLINE – FIRST($n, t_1, \dots, t_n, d_1, \dots, d_n$)

```
1: SORT jobs so that  $d_1 \leq d_2 \leq \dots \leq d_n$ .  
2:  $t \leftarrow 0$   
3: for  $j = 1$  to  $n$  do  
4:   Assign job  $j$  to interval  $[t, t + t_j]$ .  
5:    $s_j \leftarrow t; f_j \leftarrow t + t_j$   
6:    $t \leftarrow t + t_j$   
7: end for  
8: return Intervals  $[s_1, f_1], [s_2, f_2], \dots, [s_n, f_n]$ .
```




- *Observation 1.* There exists an optimal schedule with no idle time.
- *Observation 2.* The earliest-deadline-first schedule has no idle time.




Minimizing Lateness: Inversions

Def. Given a schedule S , an *inversion* is a pair of jobs i and j such that: $d_i < d_j$ but j scheduled before i .

- *Observation 3.* The earliest-deadline-first schedule has no inversions.
- *Observation 4.* If a schedule (with no idle time) has an inversion, it has a pair of inverted jobs scheduled consecutively.



Minimizing Lateness: Inversions



Claim. Swapping two adjacent, inverted jobs reduces the number of inversions by one and does not increase the max lateness.

Pf. Let l be the lateness before the swap, and let l' be it afterwards.

$$l'_k = l_k \text{ for all } k \neq i, j.$$

$$l'_i \leq l_i.$$

$$l'_j = f'_j - d_j = f_i - d_j \leq f_i - d_i \leq l_i. \quad \square$$



Analysis of Greedy Algorithm

Theorem 2

The earliest-deadline-first schedule S is optimal.

Pf. [by contradiction] Define S^* to be an optimal schedule that has the fewest number of inversions.

Can assume S^* has no idle time.

If S^* has no inversions, then $S = S^*$.

If S^* has an inversion, let $i - j$ be an adjacent inversion.


Swapping i and j : Does not increase the max lateness, but strictly decreases the number of inversions.

This contradicts definition of S^* . \square



Shortest-Paths Problem

- *Problem.* Given a digraph $G = (V, E)$, edge lengths $l_e \geq 0$, source $s \in V$, and destination $t \in V$, find the shortest directed path from s to t .



Cost of path $s-2-3-5-t$
= $9 + 23 + 2 + 16$
= 48.



Dijkstra's Algorithm



- Greedy approach. Maintain a set of explored nodes S for which algorithm has determined the shortest path distance $d(u)$ from s to u .

Initialize $S = \{s\}$, $d(s) = 0$.

Repeatedly choose unexplored node v which minimizes

$$\pi(v) = \min_{e=(u,v): u \in S} d(u) + \ell_e,$$

add v to S , and set $d(v) = \pi(v)$.



Dijkstra's Algorithm: Proof of Correctness

- *Invariant.* For each node $u \in S$, $d(u)$ is the length of the shortest $s \rightarrow u$ path.


Pf. [by induction on $|S|$]

Base case: $|S| = 1$ is easy since $S = \{s\}$ and $d(s) = 0$.

Inductive hypothesis: Assume true for $|S| = k \geq 1$.

Let v be next node added to S , and let (u, v) be the final edge.

The shortest $s \rightarrow u$ path plus (u, v) is an $s \rightarrow v$ path of length $\pi(v)$.




Dijkstra's Algorithm: Proof of Correctness—Con't

Consider any $s \rightarrow v$ path P . We show that it is no shorter than $\pi(v)$.

Let (x, y) be the first edge in P that leaves S , and let P' be the subpath to x .

P is already too long as soon as it reaches y .

$$l(P) \geq l(P') + l(x, y) \geq d(x) + l(x, y) \geq \pi(y) \geq \pi(v). \square$$



Dijkstra's Algorithm

Dijkstra's Algorithm(G, l)

- 1: Let S be the set of explored nodes.
- 2: For each $u \in S$, we store a distance $d(u)$.
- 3: Initially $S \leftarrow \{s\}$ and $d(s) \leftarrow 0$.
- 4: **while** $S \neq V$ **do**
- 5: Select a node $v \notin S$ with at least one edge from S for which $d'(v) = \min_{e=(u,v): u \in S} d(u) + l_e$ is as small as possible.
- 6: Add v to S and define $d(v) \leftarrow d'(v)$
- 7: **end while**
- 8: **return** S .

Theorem 3

Dijkstra's algorithm can find the shortest path in $O(n^2)$ time.



Homework

- Read Chapter 4 of the textbook.
- Exercises 4 & 6 in Chapter 4.

